

Fri Feb 16 15:45:42 2024time.goPage 1	Fri Feb 16 15:45:42 2024time.goPage 2
<pre>// Copyright 2009 The Go Authors. All rights reserved. // Use of this source code is governed by a BSD-style // license that can be found in the LICENSE file.  // Time-related runtime and pieces of package time.  package runtime  import (     "internal/abi"     "runtime/internal/atomic"     "runtime/internal/sys"     "unsafe" )  // A timer is a potentially repeating trigger for calling t.f(t.arg, t.seq). // Timers are allocated by client code, often as part of other data structures. // Each P has a heap of pointers to timers that it manages. // // A timer is expected to be used by only one client goroutine at a time, // but there will be concurrent access by the P managing that timer. // The fundamental state about the timer is managed in the atomic state field, // including a lock bit to manage access to the other fields. // The lock bit supports a manual cas-based spin lock that handles // contention by yielding the OS thread. The expectation is that critical // sections are very short and contention on the lock bit is low. // // Package time knows the layout of this structure. // If this struct changes, adjust ../time/sleep.go:runtimeTimer. type timer struct {     ts *timers      // Timer wakes up at when, and then at when+period, ... (period &gt; 0 only)     // each time calling f(arg, now) in the timer goroutine, so f must be     // a well-behaved function and not block.     //     // when must be positive on an active timer.     // Timers in heaps are ordered by when.     when int64     period int64     f func(any, uintptr, int64)     arg any     seq uintptr      // The state field holds state bits, defined below.     state atomic.Uintptr      // nextWhen is the next value for when,     // set if state&amp;timerNextWhen is true.     // In that case, the actual update of when = nextWhen     // must be delayed until the heap can be fixed at the same time.     nextWhen int64 }  // A timers is a per-P set of timers. type timers struct {     // lock protects timers; timers are per-P, but the scheduler can     // access the timers of another P, so we have to lock.     lock mutex      // heap is the set of timers, ordered by t.when.     // Must hold lock to access.     heap []*timer      // len is an atomic copy of len(heap).     len atomic.Uint32</pre>	<pre>// zombies is the number of timers in the heap // that are marked for removal. zombies atomic.Uint32  // raceCtx is the race context used while executing timer functions. raceCtx uintptr  // timer0When is an atomic copy of of heap[0].when. // If len(heap) == 0, timer0When is 0. timer0When atomic.Int64  // timerModifiedEarliest holds the earliest known heap[i].nextWhen field // for the heap entries with a new nextWhen pending // (that is, with the timerNextWhen bit set in t.state). // Because timers can be modified multiple times, // timerModifiedEarliest can be set to a nextWhen that has since // been replaced with a later time. // If this is 0, it means there are no timerNextWhen timers in the heap. timerModifiedEarliest atomic.Int64 }  // Timer state field. const (     // timerLocked is set when the timer is locked,     // meaning other goroutines cannot read or write mutable fields.     // Goroutines can still read the state word atomically to see     // what the state was before it was locked.     // The lock is implemented as a cas on the state field with osyield on conten     // the expectation is very short critical sections with little to no contenti     timerLocked = 1 &lt;&lt; iota      // timerHeaped is set when the timer is stored in some P's heap.     timerHeaped      // timerNextWhen is set when a pending change to the timer's when     // field has been stored in t.nextwhen. The change to t.when waits     // until the heap in which the timer appears can also be updated.     // Only set when timerHeaped is also set.     timerNextWhen      // timerZombie is set when the timer has been stopped     // but is still present in some P's heap.     // Only set when timerHeaped is also set.     // It is possible for timerNextWhen and timerZombie to both     // be set, meaning that the timer was modified and then stopped.     // A timer sending to a channel may be placed in timerZombie     // to take it out of the heap even though the timer is not stopped,     // as long as nothing is reading from the channel.     timerZombie      // timerChan is set when a timer is recognized as only existing     // to send to a channel. We take the timer out of the heap when     // nothing is watching the channel, so that the channel and     // the timer can be garbage collected if they become unreferenced,     // even if the timer is still pending.     timerChan      // timerBlocked is a value added repeatedly to the state, once per     // goroutine blocked on a timerChan timer.     // (That is, the number of blocked goroutines is state/timerBlocked.)     // Must be last, since it is not just a single bit.     timerBlocked )  // lock locks the timer, allowing reading or writing any of the timer fields.</pre>

Fri Feb 16 15:45:42 2024time.goPage 3	Fri Feb 16 15:45:42 2024time.goPage 4
<pre>// It returns the current m and the status prior to the lock. // The caller must call unlock with the same m and an updated status. func (t *timer) lock() (state uintptr, mp *m) {     acquireLockRank(lockRankTimer)     for {         state := t.state.Load()         if state&amp;timerLocked != 0 {             osyield()             continue         }         // Prevent preemption while the timer is locked.         // This could lead to a self-deadlock. See #38070.         mp := acquirem()         if t.state.CompareAndSwap(state, state timerLocked) {             return state, mp         }         releasem(mp)     } }  // unlock unlocks the timer. // If mp == nil, the caller is responsible for calling // releasem(mp) with the mp returned by t.lock. func (t *timer) unlock(state uintptr, mp *m) {     releaseLockRank(lockRankTimer)     if t.state.Load()&amp;timerLocked == 0 {         badTimer()     }     if state&amp;timerLocked != 0 {         badTimer()     }     t.state.Store(state)     if mp != nil {         releasem(mp)     } }  // initChan checks to see if a timer exists to feed a channel. // If so, it sets the timerChan bit in the state and also records // the timer in the channel's c.timer field. // initChan returns the updated state, to be passed to t.unlock. func (t *timer) initChan(state uintptr) uintptr {     if state&amp;timerChan == 0 &amp;&amp; t.arg != nil {         if e := efaceOf(&amp;t.arg); e._type.Kind() == kindChan {             state  = timerChan             t.hchan(state).timer = t         }     }     return state }  // hchan returns the channel associated with the timer t. // It must only be called for timerChan timers. func (t *timer) hchan(state uintptr) *hchan {     if state&amp;timerChan == 0 {         badTimer()     }     return (*hchan)(efaceOf(&amp;t.arg).data) }  // updateHeap updates t.when as directed by state, returning the new state // and a bool indicating whether the state (and t.when) changed. // If ts != nil, then t must be ts.heap[0], and updateHeap takes care of // moving t within the timers heap to preserve the heap invariants. // If ts == nil, then t must not be in a heap (or is in a heap that is // temporarily not maintaining its invariant, such as during timers.adjust).</pre>	<pre>func (t *timer) updateHeap(state uintptr, ts *timers) (newState uintptr, updated     if state&amp;timerZombie != 0 {         // Take timer out of heap, applying final t.when update first.         state &amp;^= timerHeaped   timerZombie         if state&amp;timerNextWhen != 0 {             state &amp;^= timerNextWhen             t.when = t.nextWhen         }         if ts != nil {             if t != ts.heap[0] {                 badTimer()             }             ts.zombies.Add(-1)             ts.deleteMin()         }         return state, true     }      if state&amp;timerNextWhen != 0 {         // Apply t.when update and move within heap.         state &amp;^= timerNextWhen         t.when = t.nextWhen         // Move t to the right position.         if ts != nil {             if t != ts.heap[0] {                 badTimer()             }             ts.siftDown(0)             ts.updateTimer0When()         }         return state, true     }      return state, false }  // maxWhen is the maximum value for timer's when field. const maxWhen = 1&lt;&lt;63 - 1  // verifyTimers can be set to true to add debugging checks that the // timer heaps are valid. const verifyTimers = false  // Package time APIs. // Godoc uses the comments in package time, not these.  // time.now is implemented in assembly.  // timeSleep puts the current goroutine to sleep for at least ns nanoseconds. // //go:linkname timeSleep time.Sleep func timeSleep(ns int64) {     if ns &lt;= 0 {         return     }      gp := getg()     t := gp.timer     if t == nil {         t = new(timer)         gp.timer = t     }     t.f = goroutineReady     t.arg = gp     t.nextWhen = nanotime() + ns     if t.nextWhen &lt; 0 { // check for overflow.</pre>

Fri Feb 16 15:45:42 2024time.goPage 5	Fri Feb 16 15:45:42 2024time.goPage 6
<pre>         t.nextWhen = maxWhen     }     gopark(resetForSleep, unsafe.Pointer(t), waitReasonSleep, traceBlockSleep, 1) }  // resetForSleep is called after the goroutine is parked for timeSleep. // We can't call resettimer in timeSleep itself because if this is a short // sleep and there are many goroutines then the P can wind up running the // timer function, goroutineReady, before the goroutine has been parked. func resetForSleep(gp *g, ut unsafe.Pointer) bool {     t := (*timer)(ut)     t.reset(t.nextWhen)     return true }  // startTimer adds t to the timer heap. // //go:linkname startTimer time.startTimer func startTimer(t *timer) {     if raceenabled {         racerelease(unsafe.Pointer(t))     }     if t.state.Load() != 0 {         throw("startTimer called with initialized timer")     }     t.reset(t.when) }  // stopTimer stops a timer. // It reports whether t was stopped before being run. // //go:linkname stopTimer time.stopTimer func stopTimer(t *timer) bool {     return t.stop() }  // resetTimer resets an inactive timer, adding it to the heap. // // Reports whether the timer was modified before it was run. // //go:linkname resetTimer time.resetTimer func resetTimer(t *timer, when int64) bool {     if raceenabled {         racerelease(unsafe.Pointer(t))     }     return t.reset(when) }  // modTimer modifies an existing timer. // //go:linkname modTimer time.modTimer func modTimer(t *timer, when, period int64) {     t.modify(when, period, t.f, t.arg, t.seq) }  // Go runtime.  // Ready the goroutine arg. func goroutineReady(arg any, _ uintptr, _ int64) {     goready(arg.(*g), 0) }  // add adds t to the timers. // The caller must have set t.ts = t, unlocked t, // and then locked ts.lock. func (ts *timers) add(t *timer) { </pre>	<pre> // Timers rely on the network poller, so make sure the poller // has started. if netpollInited.Load() == 0 {     netpollGenericInit() }  if t.ts != ts {     throw("timers.add: ts not set in timer") } ts.heap = append(ts.heap, t) ts.siftUp(len(ts.heap) - 1) if t == ts.heap[0] {     ts.updateTimer0When() } ts.len.Store(uint32(len(ts.heap))) }  // stop stops the timer t. It may be on some other P, so we can't // actually remove it from the timers heap. We can only mark it as stopped. // It will be removed in due course by the P whose heap it is on. // Reports whether the timer was stopped before it was run. func (t *timer) stop() bool {     state, mp := t.lock() Redo:     pending := false     switch {     case state&amp;timerHeaped != 0:         // Timer is in some heap, but is possibly already stopped         // (indicated by a nextWhen update to 0).         if state&amp;timerNextWhen == 0    t.nextWhen &gt; 0 {             // Timer pending: stop it.             t.nextWhen = 0             state  = timerNextWhen             pending = true         }         // Mark timer for removal unless already marked.         // (A timerChan timer might be marked for removal but not yet stopped.)         if state&amp;timerZombie == 0 {             state  = timerZombie             t.ts.zombies.Add(1)         }     case state&amp;timerChan != 0 &amp;&amp; t.when != 0:         // Active timer attached to channel but not in heap, because         // nothing is waiting on the channel or timer is stopped.         // If it should have triggered already (but nothing looked yet),         // trigger now, so that a receive after the stop sees the "old"         // value that should be there.         if state &gt;= timerBlocked { // state&amp;timerHeaped == 0             badTimer()         }     }     if now := nanotime(); t.when &lt;= now {         systemstack(func() {             t.unlockAndRun(now, state, mp) // resets t.when         })         state, mp = t.lock()         if state&amp;timerHeaped != 0 {             // While it was unlocked to run the channel send,             // the timer moved into the heap.             // Behave as though the send happened long ago             // and stop was just called now.             goto Redo         }     }     pending = t.when &gt; 0     t.when = 0 </pre>

Fri Feb 16 15:45:42 2024time.goPage 7	Fri Feb 16 15:45:42 2024time.goPage 8
<pre>     }      t.unlock(state, mp)     return pending }  // deleteMin removes timer 0 from ts. // ts must be locked. func (ts *timers) deleteMin() {     t := ts.heap[0]     if t.ts != ts {         throw("deleteMin: wrong timers")     }     t.ts = nil     last := len(ts.heap) - 1     if last &gt; 0 {         ts.heap[0] = ts.heap[last]     }     ts.heap[last] = nil     ts.heap = ts.heap[:last]     if last &gt; 0 {         ts.siftDown(0)     }     ts.updateTimer0When()     ts.len.Store(uint32(last))     if last == 0 {         // If there are no timers, then clearly none are modified.         ts.timerModifiedEarliest.Store(0)     } }  // modify modifies an existing timer. // This is called by the netpoll code or time.Ticker.Reset or time.Timer.Reset. // Reports whether the timer was modified before it was run. func (t *timer) modify(when, period int64, f func(any, uintptr, int64), arg any,     if when &lt;= 0 {         throw("timer when must be positive")     }     if period &lt; 0 {         throw("timer period must be non-negative")     }      state, mp := t.lock() Redo:     t.period = period     t.f = f     t.arg = arg     t.seq = seq      if state&amp;timerHeaped == 0 {         // Timer not in any heap, so either stopped/new         // or a timer for a currently unused channel.         // If this is a timer for a channel, initialize but leave out of heap,         // so that GC can collect it. The channel code will add the timer         // to the heap as needed to serve blocked channel ops.         // See enqueueTimerChan, dequeueTimerChan.         if state = t.initChan(state); state&amp;timerChan != 0 &amp;&amp; state &lt; timerBlocked {             pending := false             if t.when != 0 {                 if now := nanotime(); t.when &lt;= now {                     systemstack(func() {                         t.unlockAndRun(now, state, mp) // resets t.when                     })                     state, mp = t.lock()                     if state&amp;timerHeaped != 0 {                         // While it is unlocked to run the channel send, </pre>	<pre>         // the timer moved into the heap. Behave as though         // the channel send happened long ago and the         // modify call just started at this instant.         goto Redo     } }      pending = t.when &gt; 0     t.when = when     t.unlock(state, mp)     return pending }  // Not a timer for a channel, so needs to go into heap. // Assigning to when is permitted because the timer // is not in any heap, so the assignment cannot // break heap invariants. t.when = when t.unlockAndQueue(state, mp) return false }  pending := true // in the heap  if state&amp;timerZombie != 0 {     // In the heap but marked for removal.     // Therefore not pending; unmark it.     pending = false     t.ts.zombies.Add(-1)     state &amp;^= timerZombie }  // The timer is in some P's heap (perhaps another P), // so we can't change the when field. // If we did, the other P's heap would be out of order. // So we put the new when value in the nextWhen field // and set timerNextWhen, leaving the other P set the when // field when it is prepared to maintain the heap invariant. t.nextWhen = when state  = timerNextWhen earlier := when &lt; t.when if earlier {     t.ts.updateTimerModifiedEarliest(when) } t.unlock(state, mp)  // If the new status is earlier, wake up the poller. if earlier {     wakeNetPoller(when) }  return pending }  // unlockAndQueue unlocks the timer and adds it to the // local P's timer heap. func (t *timer) unlockAndQueue(state uintptr, mp *m) {     // Set up t for insertion but unlock first,     // to avoid lock inversion with timers lock.     // We set t.ts = ts so that any other concurrent     // updates to t after the unlock update the     // various atomic state in ts correctly,     // as if t were already in ts.     ts := &amp;getg().m.p.ptr().timers     state  = timerHeaped     t.ts = ts </pre>

Fri Feb 16 15:45:42 2024time.goPage 9	Fri Feb 16 15:45:42 2024time.goPage 10
<pre> when := t.when t.unlock(state, nil)  lock(&amp;ts.lock) ts.add(t) unlock(&amp;ts.lock) releasem(mp) wakeNetPoller(when) }  // reset resets the time when a timer should fire. // If used for an inactive timer, the timer will become active. // This should be called instead of addtimer if the timer value has been, // or may have been, used previously. // Reports whether the timer was active and was stopped. func (t *timer) reset(when int64) bool {     return t.modify(when, t.period, t.f, t.arg, t.seq) }  // cleanHead cleans up the head of the timer queue. This speeds up // programs that create and delete timers; leaving them in the heap // slows down addtimer. // The caller must have locked ts. func (ts *timers) cleanHead() {     gp := getg()     for {         if len(ts.heap) == 0 {             return         }          // This loop can theoretically run for a while, and because         // it is holding timersLock it cannot be preempted.         // If someone is trying to preempt us, just return.         // We can clean the timers later.         if gp.preemptStop {             return         }          t := ts.heap[0]         if t.ts != ts {             throw("timers.cleanHead: bad ts")         }          if t.state.Load()&amp;(timerNextWhen timerZombie) == 0 {             // Fast path: head of timers does not need adjustment.             return         }          state, mp := t.lock()         state, updated := t.updateHeap(state, ts)         t.unlock(state, mp)         if !updated {             // Head of timers does not need adjustment.             return         }     } }  // take moves any timers from src into ts // and then clears the timer state from src, // because src is being destroyed. // The caller must not have locked either timers. func (ts *timers) take(src *timers) {     if len(src.heap) &gt; 0 {         // The world is stopped, but we acquire timersLock to         // protect against sysmon calling timeSleepUntil. </pre>	<pre> // This is the only case where we hold the timersLock of // more than one P, so there are no deadlock concerns. lock(&amp;src.lock) lock(&amp;ts.lock) ts.move(src.heap) src.heap = nil src.len.Store(0) src.zombies.Store(0) src.timer0When.Store(0) src.timerModifiedEarliest.Store(0) unlock(&amp;ts.lock) unlock(&amp;src.lock)     } }  // moveTimers moves a slice of timers to pp. The slice has been taken // from a different P. // This is currently called when the world is stopped, but the caller // is expected to have locked the timers for pp. func (ts *timers) move(timers []*timer) {     for _, t := range timers {         state, mp := t.lock()         t.ts = nil         state, _ = t.updateHeap(state, nil)         // Unlock before add, to avoid append (allocation)         // while holding lock. This would be correct even if the world wasn't         // stopped (but it is), and it makes staticclockranking happy.         if state&amp;timerHeaped != 0 {             t.ts = ts         }         t.unlock(state, mp)         if state&amp;timerHeaped != 0 {             ts.add(t)         }     } }  // adjust looks through the timers in the current P's heap for // any timers that have been modified to run earlier, and puts them in // the correct place in the heap. While looking for those timers, // it also moves timers that have been modified to run later, // and removes deleted timers. The caller must have locked the timers for pp. func (ts *timers) adjust(now int64, force bool) {     // If we haven't yet reached the time of the earliest timerModified     // timer, don't do anything. This speeds up programs that adjust     // a lot of timers back and forth if the timers rarely expire.     // We'll postpone looking through all the adjusted timers until     // one would actually expire.     if !force {         first := ts.timerModifiedEarliest.Load()         if first == 0    first &gt; now {             if verifyTimers {                 ts.verify()             }             return         }     }      // We are going to clear all timerModified timers.     ts.timerModifiedEarliest.Store(0)      changed := false     for i := 0; i &lt; len(ts.heap); i++ {         t := ts.heap[i]         if t.ts != ts {             throw("timers.adjust: bad ts") </pre>

Fri Feb 16 15:45:42 2024time.goPage 11	Fri Feb 16 15:45:42 2024time.goPage 12
<pre>     }      state, mp := t.lock()     if state&amp;timerHeaped == 0 {         badTimer()     }     state, updated := t.updateHeap(state, nil)     if updated {         changed = true         if state&amp;timerHeaped == 0 {             n := len(ts.heap)             ts.heap[i] = ts.heap[n-1]             ts.heap[n-1] = nil             ts.heap = ts.heap[:n-1]             t.ts = nil             ts.zombies.Add(-1)             i--         }     }     t.unlock(state, mp) }  if changed {     ts.initHeap()     ts.updateTimer0When() }  if verifyTimers {     ts.verify() } }  // wakeTime looks at ts's timers and returns the time when we // should wake up the netpoller. It returns 0 if there are no timers. // This function is invoked when dropping a P, so it must run without // any write barriers. // //go:nowritebarrierrec func (ts *timers) wakeTime() int64 {     next := ts.timer0When.Load()     nextAdj := ts.timerModifiedEarliest.Load()     if next == 0    (nextAdj != 0 &amp;&amp; nextAdj &lt; next) {         next = nextAdj     }     return next }  // check runs any timers for the P that are ready. // If now is not 0 it is the current time. // It returns the passed time or the current time if now was passed as 0. // and the time when the next timer should run or 0 if there is no next timer, // and reports whether it ran any timers. // If the time when the next timer should run is not 0, // it is always larger than the returned time. // We pass now in and out to avoid extra calls of nanotime. // //go:yeswritebarrierrec func (ts *timers) check(now int64) (rnow, pollUntil int64, ran bool) {     // If it's not yet time for the first timer, or the first adjusted     // timer, then there is nothing to do.     next := ts.wakeTime()     if next == 0 {         // No timers to run or adjust.         return now, 0, false     } </pre>	<pre>     if now == 0 {         now = nanotime()     }      // If this is the local P, and there are a lot of deleted timers,     // clear them out. We only do this for the local P to reduce     // lock contention on timersLock.     force := ts == &amp;getg().m.p.ptr().timers &amp;&amp; int(ts.zombies.Load()) &gt; int(ts.le      if now &lt; next &amp;&amp; !force {         // Next timer is not ready to run, and we don't need to clear deleted tim         return now, next, false     }      lock(&amp;ts.lock)     if len(ts.heap) &gt; 0 {         ts.adjust(now, force)         for len(ts.heap) &gt; 0 {             // Note that runtime may temporarily unlock             // pp.timersLock.             if tw := ts.run(now); tw != 0 {                 if tw &gt; 0 {                     pollUntil = tw                 }                 break             }             ran = true         }     }      unlock(&amp;ts.lock)      return now, pollUntil, ran }  // run examines the first timer in timers. If it is ready based on now, // it runs the timer and removes or updates it. // Returns 0 if it ran a timer, -1 if there are no more timers, or the time // when the first timer should run. // The caller must have locked the timers for pp. // If a timer is run, this will temporarily unlock the timers. // //go:systemstack func (ts *timers) run(now int64) int64 { Redo:     if len(ts.heap) == 0 {         return -1     }     t := ts.heap[0]     if t.ts != ts {         throw("timers.run: bad ts")     }      if t.state.Load()&amp;(timerNextWhen timerZombie) == 0 &amp;&amp; t.when &gt; now {         // Fast path: not ready to run.         // The access of t.when is protected by the caller holding         // pp.timersLock, even though t itself is unlocked.         return t.when     }      state, mp := t.lock()     state, updated := t.updateHeap(state, ts)     if updated {         t.unlock(state, mp)         goto Redo     } </pre>

Fri Feb 16 15:45:42 2024time.goPage 13	Fri Feb 16 15:45:42 2024time.goPage 14
<pre>     if state&amp;timerHeaped == 0 {         badTimer()     }      if t.when &gt; now {         // Not ready to run.         t.unlock(state, mp)         return t.when     }      t.unlockAndRun(now, state, mp)     return 0 }  // unlockAndRun unlocks and runs the timer t. // If t is in a timer set (t.ts != nil), the caller must have locked the timer set // and this call will temporarily unlock the timer set while running the timer func // //go:systemstack func (t *timer) unlockAndRun(now int64, state uintptr, mp *m) {     if raceenabled {         // Note that we are running on a system stack,         // so there is no chance of getg().m being reassigned         // out from under us while this function executes.         tsLocal := &amp;getg().m.p.ptr().timers         if tsLocal.raceCtx == 0 {             tsLocal.raceCtx = racegostart(abi.FuncPCABIInternal((*timers).run) +             raceacquirectx(tsLocal.raceCtx, unsafe.Pointer(t)))         }     }      if state&amp;(timerNextWhen timerZombie) != 0 {         badTimer()     }      f := t.f     arg := t.arg     seq := t.seq     var next int64     delay := now - t.when     if t.period &gt; 0 {         // Leave in heap but adjust next time to fire.         next = t.when + t.period*(1+delay/t.period)         if next &lt; 0 { // check for overflow.             next = maxWhen         }     } else {         next = 0     }     if state&amp;timerHeaped != 0 {         t.nextWhen = next         state  = timerNextWhen         if next == 0 {             state  = timerZombie         }     } else {         t.when = next     }     ts := t.ts     state, _ = t.updateHeap(state, ts)     t.unlock(state, mp)      if raceenabled {         // Temporarily use the current P's racectx for g0.         gp := getg() </pre>	<pre>         if gp.racectx != 0 {             throw("timers.run: unexpected racectx")         }         gp.racectx = gp.m.p.ptr().timers.raceCtx     }      if ts != nil {         unlock(&amp;ts.lock)     }     f(arg, seq, delay)     if ts != nil {         lock(&amp;ts.lock)     }      if raceenabled {         gp := getg()         gp.racectx = 0     } }  // updateTimerPMask clears pp's timer mask if it has no timers on its heap. // // Ideally, the timer mask would be kept immediately consistent on any timer // operations. Unfortunately, updating a shared global data structure in the // timer hot path adds too much overhead in applications frequently switching // between no timers and some timers. // // As a compromise, the timer mask is updated only on pidleget / pidleput. A // running P (returned by pidleget) may add a timer at any time, so its mask // must be set. An idle P (passed to pidleput) cannot add new timers while // idle, so if it has no timers at that time, its mask may be cleared. // // Thus, we get the following effects on timer-stealing in findrunnable: // // - Idle Ps with no timers when they go idle are never checked in findrunnable //   (for work- or timer-stealing; this is the ideal case). // - Running Ps must always be checked. // - Idle Ps whose timers are stolen must continue to be checked until they run //   again, even after timer expiration. // // When the P starts running again, the mask should be set, as a timer may be // added at any time. // // TODO(prattmic): Additional targeted updates may improve the above cases. // e.g., updating the mask when stealing a timer. func updateTimerPMask(pp *p) {     if pp.timers.len.Load() &gt; 0 {         return     }      // Looks like there are no timers, however another P may transiently     // decrement numTimers when handling a timerModified timer in     // checkTimers. We must take timersLock to serialize with these changes.     lock(&amp;pp.timers.lock)     if pp.timers.len.Load() == 0 {         timerpMask.clear(pp.id)     }     unlock(&amp;pp.timers.lock) }  // verifyTimerHeap verifies that the timers is in a valid state. // This is only for debugging, and is only called if verifyTimers is true. // The caller must have locked the timers. func (ts *timers) verify() {     for i, t := range ts.heap {         if i == 0 { </pre>

Fri Feb 16 15:45:42 2024time.goPage 15	Fri Feb 16 15:45:42 2024time.goPage 16
<pre>         // First timer has no parent.         continue     }      // The heap is 4-ary. See siftupTimer and siftdownTimer.     p := (i - 1) / 4     if t.when &lt; ts.heap[p].when {         print("bad timer heap at ", i, ": ", p, ": ", ts.heap[p].when, ", ",             throw("bad timer heap")     } } if n := int(ts.len.Load()); len(ts.heap) != n {     println("timer heap len", len(ts.heap), "!= atomic len", n)     throw("bad timer heap len") }  // updateTimer0When sets the P's timer0When field. // The caller must have locked the timers for pp. func (ts *timers) updateTimer0When() {     if len(ts.heap) == 0 {         ts.timer0When.Store(0)     } else {         ts.timer0When.Store(ts.heap[0].when)     } }  // updateTimerModifiedEarliest updates the recorded nextwhen field of the // earlier timerModifiedEarlier value. // The timers for pp will not be locked. func (ts *timers) updateTimerModifiedEarliest(nextwhen int64) {     // The low bit of timerModifiedEarliest tracks how the value was set.     // Low bit 1 means it was set by updateTimerModifiedEarliest     // (without holding ts.lock).     nextwhen  = 1     for {         old := ts.timerModifiedEarliest.Load()         if old != 0 &amp;&amp; old &lt; nextwhen {             return         }          if ts.timerModifiedEarliest.CompareAndSwap(old, nextwhen) {             return         }     } }  // timeSleepUntil returns the time when the next timer should fire. Returns // maxWhen if there are no timers. // This is only called by sysmon and checkdead. func timeSleepUntil() int64 {     next := int64(maxWhen)      // Prevent allp slice changes. This is like retake.     lock(&amp;allpLock)     for _, pp := range allp {         if pp == nil {             // This can happen if procrsize has grown             // allp but not yet created new Ps.             continue         }          if w := pp.timers.wakeTime(); w != 0 {             next = min(next, w)         }     } } </pre>	<pre>         unlock(&amp;allpLock)          return next     }      // Heap maintenance algorithms.     // These algorithms check for slice index errors manually.     // Slice index error can happen if the program is using racy     // access to timers. We don't want to panic here, because     // it will cause the program to crash with a mysterious     // "panic holding locks" message. Instead, we panic while not     // holding a lock.      // siftUp puts the timer at position i in the right place     // in the heap by moving it up toward the top of the heap.     func (ts *timers) siftUp(i int) {         t := ts.heap         if i &gt;= len(t) {             badTimer()         }         when := t[i].when         if when &lt;= 0 {             badTimer()         }         tmp := t[i]         for i &gt; 0 {             p := (i - 1) / 4 // parent             if when &gt;= t[p].when {                 break             }             t[i] = t[p]             i = p         }         if tmp != t[i] {             t[i] = tmp         }     }      // siftDown puts the timer at position i in the right place     // in the heap by moving it down toward the bottom of the heap.     func (ts *timers) siftDown(i int) {         t := ts.heap         n := len(t)         if i &gt;= n {             badTimer()         }         when := t[i].when         if when &lt;= 0 {             badTimer()         }         tmp := t[i]         for {             c := i*4 + 1 // left child             c3 := c + 2 // mid child             if c &gt;= n {                 break             }             w := t[c].when             if c+1 &lt; n &amp;&amp; t[c+1].when &lt; w {                 w = t[c+1].when                 c++             }             if c3 &lt; n {                 w3 := t[c3].when                 if c3+1 &lt; n &amp;&amp; t[c3+1].when &lt; w3 {                     w3 = t[c3+1].when </pre>



Fri Feb 16 15:45:42 2024time.goPage 17	Fri Feb 16 15:45:42 2024time.goPage 18
<pre>         c3++     }     if w3 &lt; w {         w = w3         c = c3     } } if w &gt;= when {     break } t[i] = t[c] i = c } if tmp != t[i] {     t[i] = tmp } }  // initHeap reestablishes the heap order in the slice ts.heap. // It takes O(n) time for n=len(ts.heap), not the O(n log n) of n repeated add op func (ts *timers) initHeap() {     // Last possible element that needs sifting down is parent of last element;     // last element is len(t)-1; parent of last element is (len(t)-1-1)/4.     if len(ts.heap) &lt;= 1 {         return     }     for i := (len(ts.heap) - 1 - 1) / 4; i &gt;= 0; i-- {         ts.siftDown(i)     } }  // badTimer is called if the timer data structures have been corrupted, // presumably due to racy use by the program. We panic here rather than // panicking due to invalid slice access while holding locks. // See issue #25686. func badTimer() {     throw("timer data corruption") }  // Timer channels.  // maybeRunChan checks whether the timer needs to run // to send a value to its associated channel. If so, it does. // The timer must not be locked. func (t *timer) maybeRunChan() {     if t.state.Load()&amp;timerHeaped != 0 {         // If the timer is in the heap, the ordinary timer code         // is in charge of sending when appropriate.         return     }      state, mp := t.lock()     now := nanotime()     if state&amp;timerHeaped != 0    t.when == 0    t.when &gt; now {         // Timer in the heap, or not running at all, or not triggered.         t.unlock(state, mp)         return     }     systemstack(func() {         t.unlockAndRun(now, state, mp)     }) }  // enqueueTimerChan is called when a channel op has decided to block on c. // The caller holds the channel lock for c and possibly other channels. // enqueueTimerChan makes sure that c is in the timer heap, </pre>	<pre> // adding it if needed. func enqueueTimerChan(c *hchan) {     t := c.timer     state, mp := t.lock()     if state&amp;timerChan == 0 {         state = t.initChan(state)         if state&amp;timerChan == 0 {             badTimer()         }     }     state += timerBlocked     if state &gt;= 2*timerBlocked {         // Already blocked and therefore in heap if running.         if t.when &gt; 0 &amp;&amp; state&amp;timerHeaped == 0 {             badTimer()         }     }     t.unlock(state, mp)     return }  if state&amp;timerHeaped != 0 {     // Already in heap, but if this the first enqueue after a recent dequeue,     // it may be marked for removal. Unmark it if so, but don't unmark     // if the removal is because the timer is not running at all.     if state&amp;timerNextWhen == 0    t.nextWhen != 0 {         state ^= timerZombie         t.ts.zombies.Add(-1)     }     t.unlock(state, mp)     return }  if t.when == 0 {     // Timer not running. Skip adding to heap.     t.unlock(state, mp)     return }  // Not in heap, but timer is running. Need to add to heap now. t.unlockAndQueue(state, mp) }  // dequeueTimerChan is called when a channel op that was blocked on c // is no longer blocked. Every call to enqueueTimerChan must be paired with // a call to dequeueTimerChan. // The caller holds the channel lock for c and possibly other channels. // dequeueTimerChan removes c from the timer heap when nothing is // blocked on it anymore. func dequeueTimerChan(c *hchan) {     t := c.timer     state, mp := t.lock()     if state&amp;timerChan == 0    state &lt; timerBlocked {         badTimer()     }     state -= timerBlocked     if state &lt; timerBlocked &amp;&amp; state&amp;timerHeaped != 0 &amp;&amp; state&amp;timerZombie == 0 {         // Last goroutine that was blocked on this timer.         // Mark for removal from heap but do not clear t.when,         // so that we know what time it is still meant to trigger.         state  = timerZombie         t.ts.zombies.Add(1)     }     t.unlock(state, mp) } </pre>